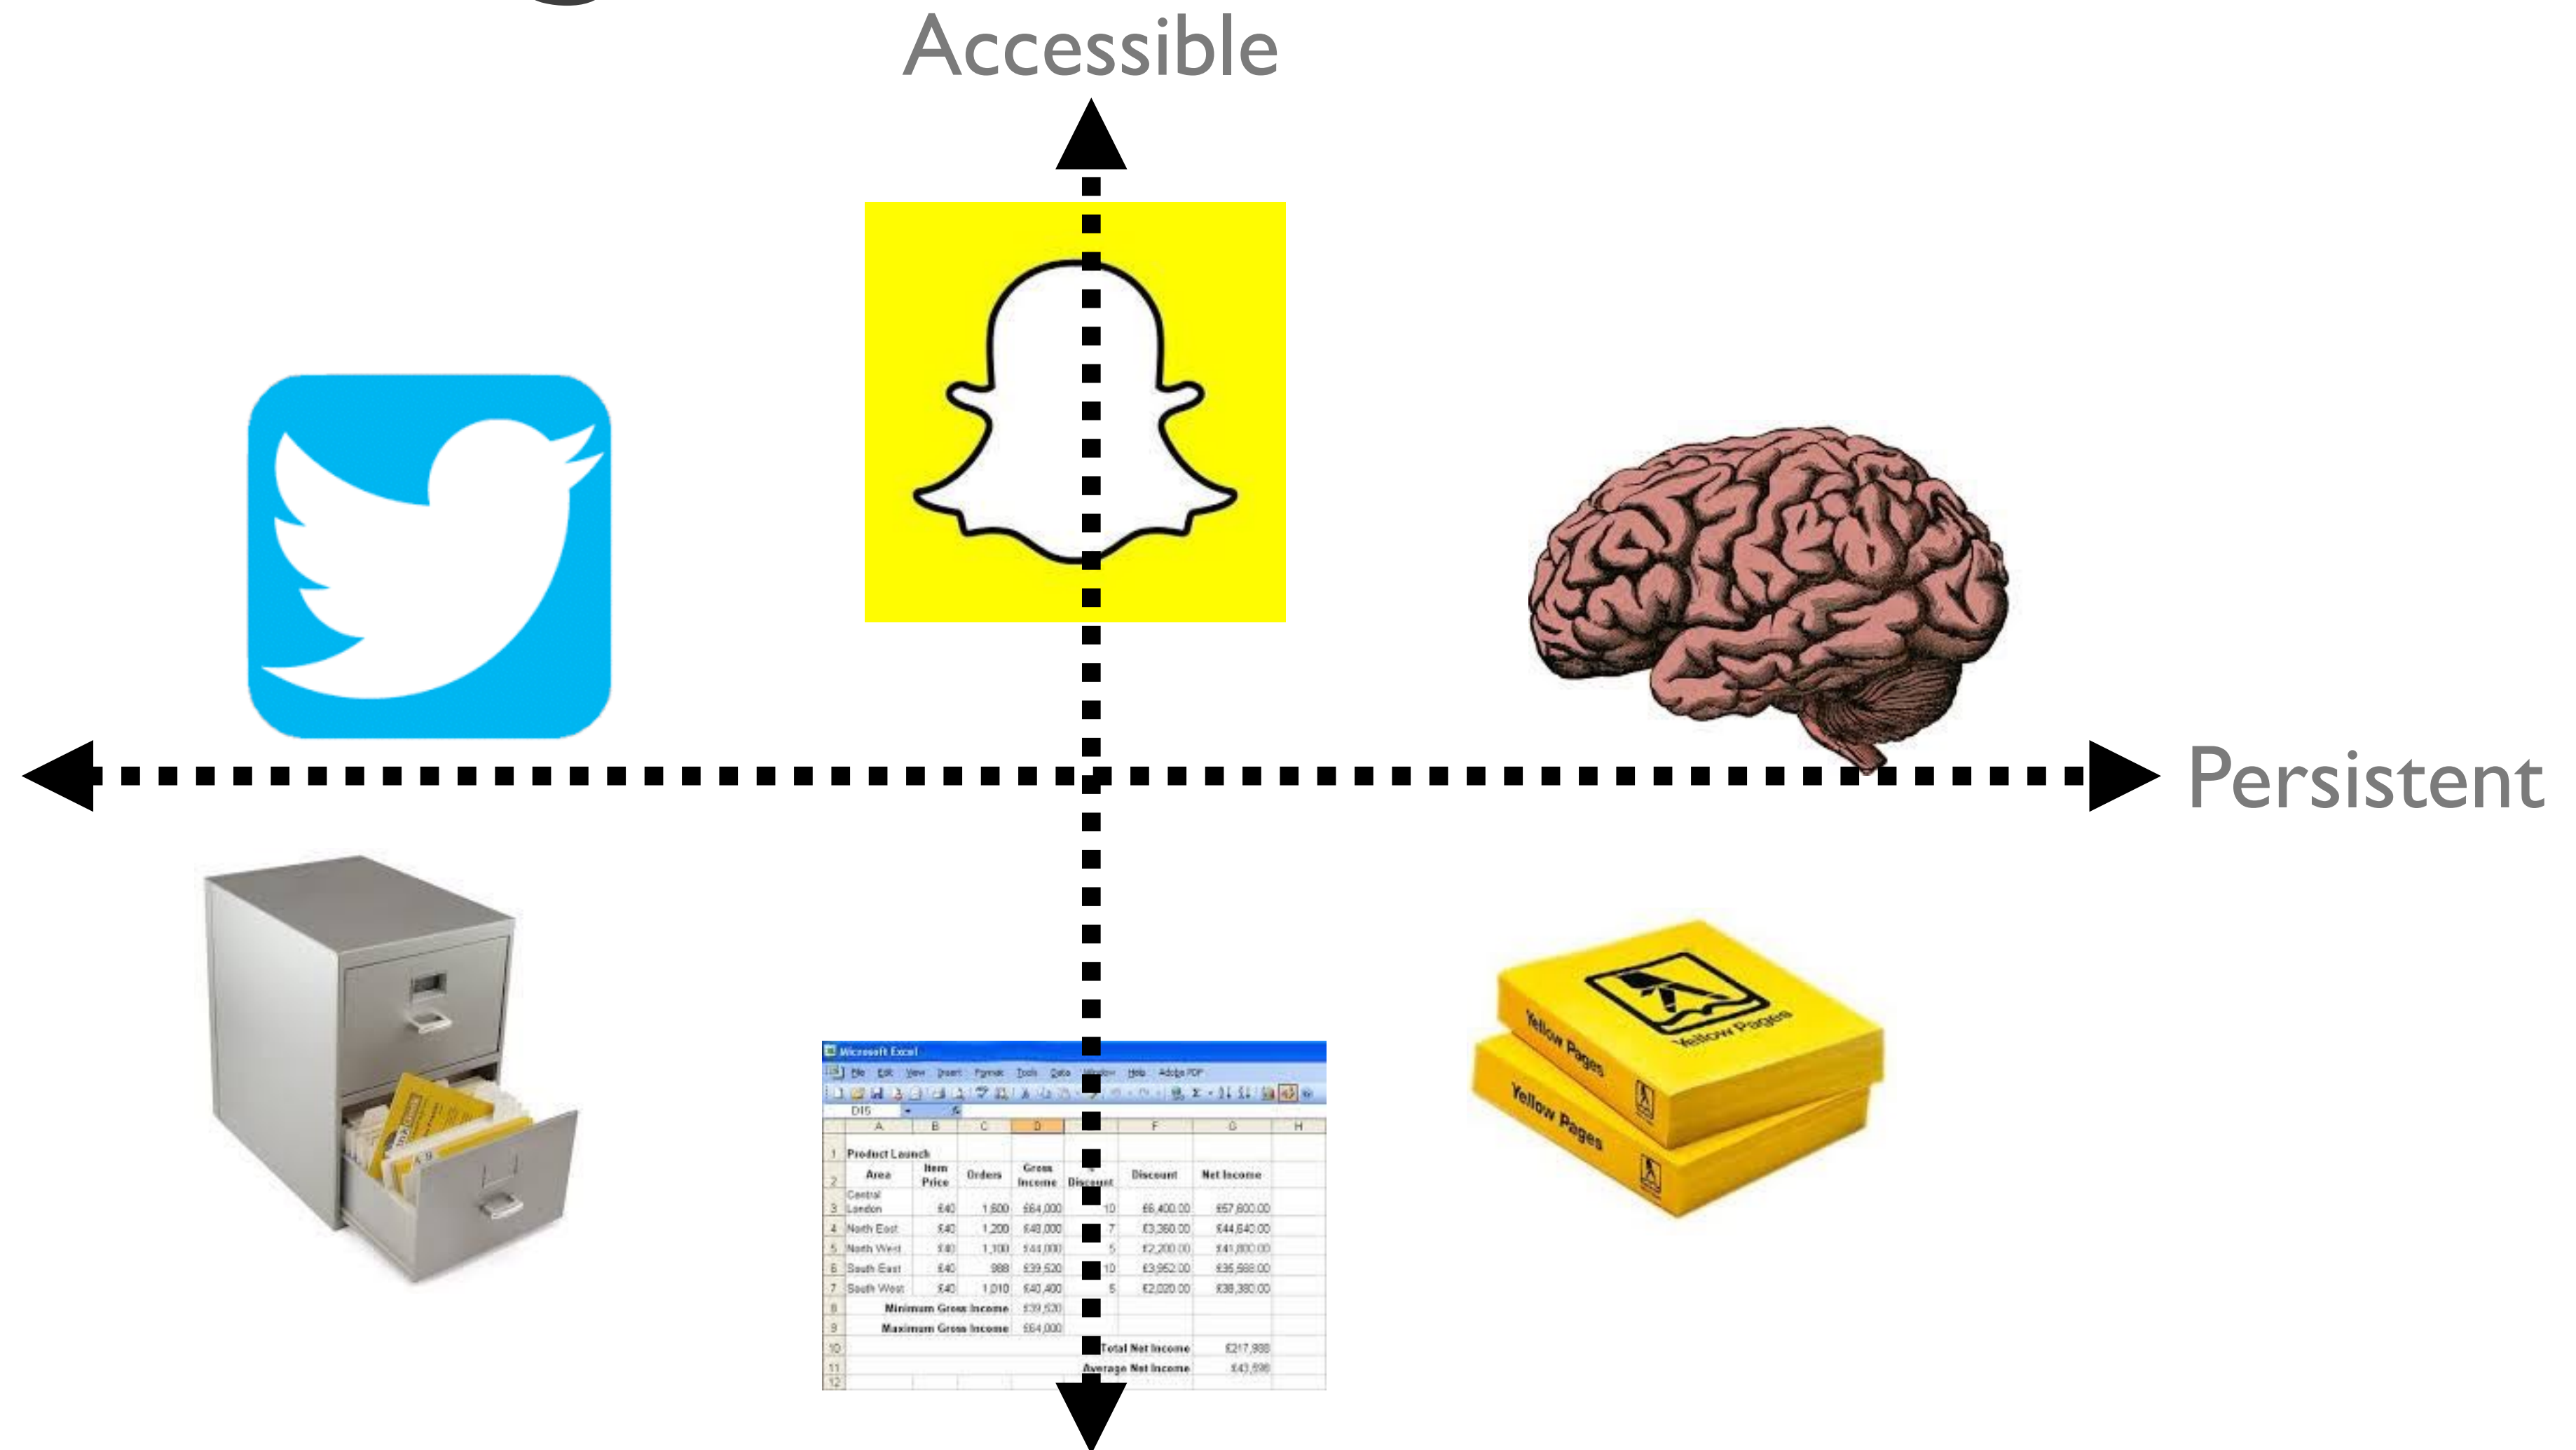


Intro to Databases

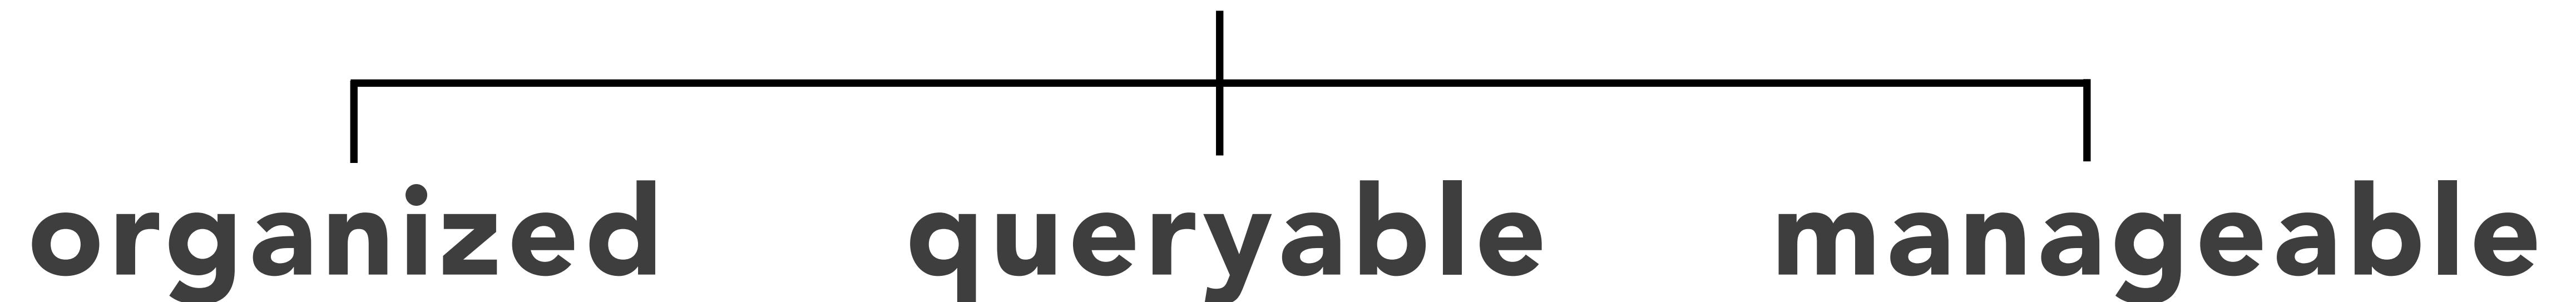
SQL

What is a database?

Things that hold info



A database **persists** information
and is **accessible** via code



Organized: Standard Storage Formatting

- DBs are a collection of **Tables** (or *relations*)
- Tables have **Columns** (*attributes* / *fields*) that describe **Rows** (*instances* / *tuples*)
- Duplicate rows are not allowed
- Rows often have a **primary key** (unique identifier)

Table / Relation

Column / Attribute / Field Column / Attribute / Field Column / Attribute / Field

	ID	Name	Type
Row / Tuple / Instance	1	Pikachu	lightning
Row / Tuple / Instance	2	Squirtle	water
Row / Tuple / Instance	3	Charmander	fire
Row / Tuple / Instance	4	Bulbasaur	grass

Queryable: via a Standard Language

- A simple, structured query language: SQL
- Declarative (vs. imperative)
- No more hand-rolled algorithms / data structures
- DBMS picks an efficient execution strategy based on indexes, data, workload etc.



SQL

```
-- Pikachu, I choose you!  
SELECT id, name  
FROM pokemon  
WHERE type = 'lightning'  
LIMIT 1
```



Manageable: Easy, Safe, Performant

- Offloads work and requisite understanding of programming
- Knowledge is portable
- Abstraction
- Transfer data between systems
- DBMS can make certain guarantees
 - prevent unsafe operations
 - built-in redundancies
 - handle multiple users, threads


How Did We End Up Here?

Before Relational DBs (ca. < 1970s)

- Data stored in custom “data files”
- Queried via application-specific code
- Advantages
 - Middle layer not needed
 - Solutions customized for each application
- Disadvantages
 - Hard to change the system
 - Knowledge not compounding
 - Data-transfer is difficult



Database Management Systems (DBMS)




- One layer and language to store and access data
- Sold as a way for “non-technical people” to manage data

“Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation).”

– E. F. CODD,
A RELATIONAL MODEL OF DATA FOR
LARGE SHARED DATA BANKS

Relational Databases & Logic



- 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- Wrote Alpha (never implemented) as a *query language*
- IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created **Structured English Query Lang**
- **SEQUEL** way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- **SQL became the standard (ANSI 1986, ISO 1987)**
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*

Appreciating Databases

- **Ubiquitous**
- **Standardized**
- **Complex / deep**
- **Powerful: database admins are**
 - Feared by developers
 - ...but also taken for granted until things break
 - Befriended by business people
 - Contacted by the government for secret data (e.g. NSA)

ACID Guarantees

- **Atomicity**
- **Consistency**
- **Isolation**
- **Durability**

What happens if the database crashes?

- Suppose we have a banking system with client accounts
- Every debit must have a matching credit.
- Imagine a crash results in only one table being updated.
 - Database inconsistency → unexpected data and software crashes/bugs
 - Financial risk for clients

Atomic Transactions

- **atomic transaction: A set of database operations that must occur together**
 - i.e. A debit to one bank account, and a credit to another
- **A transaction must either succeed or fail; it cannot partially complete.**
- **Every database query is represented by a transaction**

Consistency

- **Specify rules that columns need to follow**
 - Gender column can only contain M, F, or U.
 - Savings account must start with S or checking with C
 - Column cannot be null
- **Protect the database from inconsistencies and simplify software logic**
 - Allows software to make assumptions about underlying data

Resource Management

- Processes can be readers and writers
- Files can have many readers
- If a process has a writer, no other process can read from it, and no other process can write to it

Proposed File Scheme

- Suppose that we have decided not to use a database and instead store our data in a series of files.
- How might our setup fail to serve queries from multiple users?

Deadlock



Databases give us concurrency (Isolation)

- Multiple clients can make queries to read and update without the risk of deadlock or starvation.

Persistence/Durability

- Files are also persistence (store information without power)

Progression of Databases

- **Navigational (< 1970s)**
 - More common during tape era; entries had references to next entries.
- **Relational (> 1970s)**
 - Based on relational (table-based) logic, see E.F. Codd.
- **NoSQL (> 2000s)**
 - "Not only SQL" — document storage, for example.

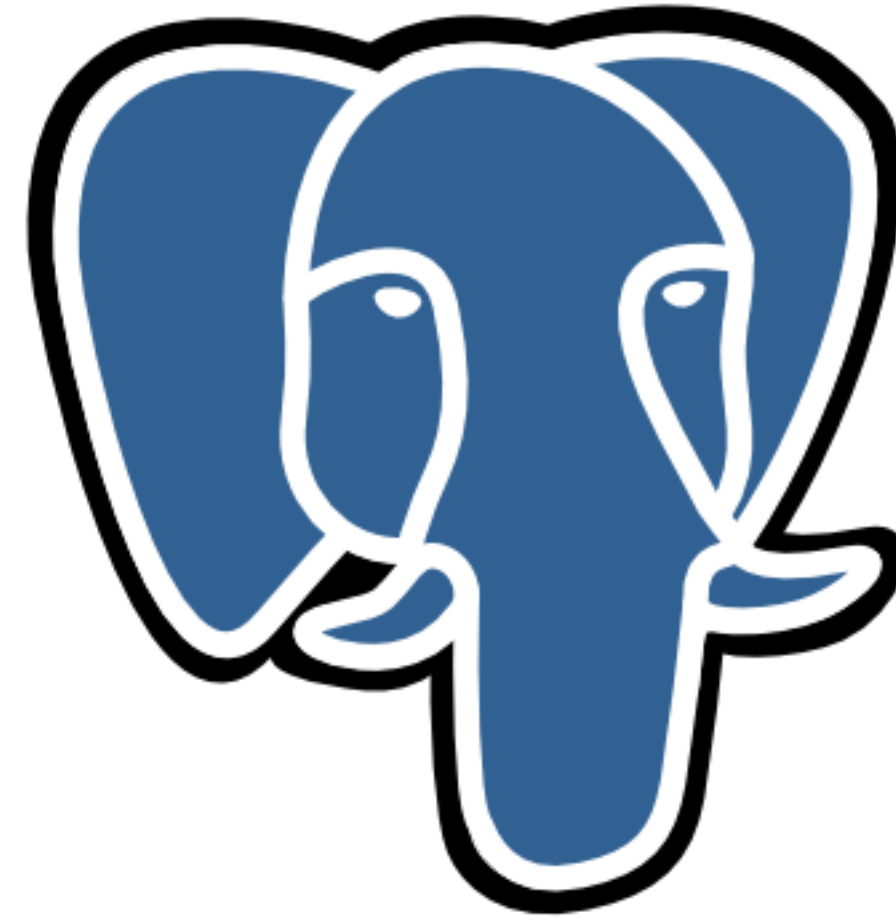
RDBMS vs NoSQL

- **A DBMS doesn't *have* to be relational**
 - Remember, DBMS is just an application that intelligently stores data and can answer requests to manage that data
- **Lately, many "NoSQL" or non-relational DBMSs have been gaining popularity**
 - Graph databases (e.g. GraphQL)
 - Document databases (e.g. MongoDB)
 - Hybrids (e.g. PostgreSQL)
- **RDBMSs still remain the #1 DB option for now**



Some well-known rDBMSs





PostgreSQL

Why PostgreSQL?

- Advanced, powerful, and popular
- Rapid open source development
- Highly extensible (stored procedures)
- Deep SQL standards compliance
- NoSQL ("Not Only SQL"), objective support
- Excellent transactions / ACID reliability; focus on integrity
- Multi-user management / administration

History of PostgreSQL

- 1970s at UC Berkeley:
INteractive **G**raphics **R**etrieval **S**ystem (INGRES)
- 1980s: POSTGRES ("Post-Ingres")
- 1995: POSTQUEL and Postgres95.
 - `monitor -> psql`
- 1996: Adopted by the open source community
 - Ongoing: stability, testing, documentation, new features
 - PostgreSQL

psql

```
1. psql (psql)
X: .../hello-world (zsh) 1 X: .../class-libs (zsh) 2 X: psql (psql) 3 X: postgres (postgres) 4
psql [local]:5432 glebec # ~ \l
List of databases

```

Name	Owner	Encoding	Collate	Ctype	Access privileges
assessmentexpresssequeliza	fullstack	UTF8	en_US.UTF-8	en_US.UTF-8	
author	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
checkpoint_angular	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
checkpoint_express_review	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
glebec	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
juke	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
postgres	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
sequelizecheckpoint	fullstack	UTF8	en_US.UTF-8	en_US.UTF-8	
template0	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	=c/glebec glebec=CTc/glebec
template1	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	=c/glebec glebec=CTc/glebec
tripolanner	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
twitterdb	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
wikistack	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	

```
(13 rows)

psql [local]:5432 glebec # ~ \c author
You are now connected to database "author" as user "glebec".

psql [local]:5432 glebec # author \dt
List of relations

```

Schem	Name	Type	Owner
public	stories	table	glebec
public	users	table	glebec

```
(2 rows)

psql [local]:5432 glebec # author
```



pgcli

```
stayupdated_test> \d
+-----+-----+-----+-----+
| Schema | Name                | Type      | Owner  |
+-----+-----+-----+-----+
| public | admins              | table     | amjith |
| public | cpes                | table     | amjith |
| public | goose_db_version    | table     | amjith |
| public | goose_db_version_id_seq | sequence  | amjith |
| public | packages            | table     | amjith |
| public | packages_id_seq     | sequence  | amjith |
| public | users               | table     | amjith |
| public | users_id_seq        | sequence  | amjith |
| public | vulnerabilities     | table     | amjith |
| public | vulnerabilities_cpes | table     | amjith |
| public | vulnerabilities_id_seq | sequence  | amjith |
+-----+-----+-----+-----+
SELECT 11
stayupdated_test> SELECT * FROM users;
+-----+-----+-----+-----+-----+
| id | display_name | password | email           | created_on |
+-----+-----+-----+-----+-----+
| 177 | DisplayName1 | 1024cms  | user@ex.com     | 2014-11-15 15:02:50.094560 |
| 180 | testname2    | pas5w0rd | email@ex.com    | 2014-11-28 10:25:46.170660 |
| 181 | amjith       | password | amjith@amjith.amjith | 2014-11-28 18:39:48.195067 |
+-----+-----+-----+-----+-----+
SELECT 3
stayupdated_test> SELECT * FROM
|
| admins
| cpes
| goose_db_version
| packages
| users
```


Postico

The screenshot shows the Postico application window. The title bar includes window controls, navigation arrows, a breadcrumb trail (Reporting > reporting > forex), a status bar (Connected.), and the database version (PostgreSQL 9.4.5). The main area is divided into a left sidebar, a central table, and a right sidebar.

Left Sidebar: A list of database tables with icons. The 'forex' table is currently selected and highlighted.

Central Table: A table with 6 columns: currency, base, rate, date, and source_id. It displays 16 rows of data for the date 2014-07-25. The first three rows are highlighted in blue.

currency	base	rate	date	source_id
AED	USD	3.67291	2014-07-25	1
AFN	USD	56.485726	2014-07-25	1
ALL	USD	103.5838	2014-07-25	1
AMD	USD	410.086	2014-07-25	1
ANG	USD	1.787	2014-07-25	1
AOA	USD	96.952626	2014-07-25	1
ARS	USD	8.169642	2014-07-25	1
AUD	USD	1.062844	2014-07-25	1
AWG	USD	1.79	2014-07-25	1
AZN	USD	0.784067	2014-07-25	1
BAM	USD	1.453024	2014-07-25	1
BBD	USD	2	2014-07-25	1
BDT	USD	77.61971	2014-07-25	1
BGN	USD	1.452543	2014-07-25	1

Right Sidebar: A panel for filtering the table data. It contains dropdown menus for 'currency' (set to 'MULTIPLE'), 'base' (set to 'USD'), 'rate' (set to 'MULTIPLE'), and 'date' (set to '2014-07-25'). Below these is a 'source_id' dropdown with a table preview showing one entry with id '1' and name 'Open Exchange R...'. At the bottom are tabs for 'Content' and 'Structure', a '+' button, a 'Filter' button, and pagination controls showing 'Page 1 of 342'.

LAB